Vladislav Fedorov

+1 (217) 974-9808 | vldslv.valeryevich@gmail.com

GitHub: github.com/ValDevelopment | LinkedIn: https://www.linkedin.com/in/fedorov-vladislav-v

Work Authorization: 36 Months STEM OPT (No Visa Sponsorship Required) | Available Immediately | Open to Relocation

PROFILE

MS Statistics graduate with experience developing reproducible analyses in R, SAS, Python, and SQL across biomedical, operational, and time-series applications. Skilled in regression, hypothesis testing, survival analysis, power analysis, and statistical programming. Seeking statistician, biostatistics, and statistical programming roles.

EDUCATION

University of Illinois Urbana-Champaign | Champaign, IL | Aug 2024 - May 2026
Master of Science in Statistics | GPA: 3.90

University of Illinois Urbana-Champaign | Champaign, IL | Aug 2020 - Dec 2023
Bachelor of Liberal Arts and Sciences in Statistics, Minor in Mathematics | GPA: 3.77; Dean's List, Fall 2023

SKILLS

Statistical Programming: R, SAS, Python (Pandas, NumPy, scikit-learn), SQL
Statistical Methods: Regression, Hypothesis Testing, Survival Analysis, Power Analysis, Time Series, Classification, Disproportionality Analysis
Tools & Technologies: Excel, MySQL, Git, GitHub, Tableau, Power BI, Streamlit
Languages: Russian (Native), English (Fluent), Italian (Intermediate)

PROFESSIONAL EXPERIENCE

Statistical Consultant | Champaign, IL
University of Illinois Urbana-Champaign | Jan 2026 - May 2026

- Conducted exploratory analysis of 6M+ job records and 100+ operational variables in Python and R, identifying resource inefficiency as a high-impact area and defining the scope of the subsequent analysis
- Built reusable Python and R programs that automate data preparation, resource-efficiency calculations, and visualizations, enabling the client to reproduce and extend the analysis with future job data
- Quantified systematic resource over-requesting, finding that the median job ran for 1.2 minutes against a 240-minute request and that approximately 800M reserved CPU-hours went unused
- Compared requested versus actual resource usage, revealing 1.48% median memory efficiency, 91.5% of jobs using less than half their requested memory, and pronounced CPU underutilization among jobs requesting 8–32 cores
- Collaborated with teammates and client stakeholders to present findings and propose personalized usage reports benchmarking each user's resource efficiency against cluster-wide patterns; the client identified the concept as a strong direction for future development

PROJECTS

FAERS Pharmacovigilance Signal Detection

- Analyzed 385,288 FDA FAERS reports, including 3,279 semaglutide cases, to identify potential adverse-event signals
- Built parallel disproportionality-analysis pipelines in Python and SAS, calculating PRR, ROR, and chi-square statistics with Haldane-Anscombe correction and reproducing identical results across both implementations
- Identified 15 adverse-event signals and evaluated them against FDA labeling and peer-reviewed literature, visualizing disproportionality metrics in an interactive Tableau dashboard

Competitive Esports Predictive Modeling

- Built logistic-regression and XGBoost win-prediction models from professional League of Legends data, achieving 0.84 AUC on a chronologically held-out season and identifying early-game objectives with limited incremental predictive value
- Forecasted weekly game-decisiveness trends using ARIMA, ETS, and gradient boosting with rolling-origin backtesting
- Built a sequential Elo rating system across 26,000+ professional matches to compare regional team strength, tune the K-factor empirically, and quantify uncertainty through bootstrapped confidence intervals

OpenPowerlifting Relational Database

- Built a reproducible Python preprocessing pipeline to clean raw competition data, generate entity identifiers, preserve missing and failed attempts, and export relational tables for database loading
- Designed a normalized MySQL database for 100,000 OpenPowerlifting competition results, separating lifter, meet, and result data into linked tables with primary and foreign key constraints
- Wrote SQL analyses using joins, conditional aggregation, CTEs, and window functions, then built interactive Power BI dashboards to analyze federation participation, performance differences, sanctioned meets, and lift-attempt success rates

Natural Language SQL Analytics Agent

- Built an LLM-powered analytics agent using Python, MySQL, Groq API, Pydantic, and SQLGlot to translate natural-language questions into schema-aware SQL queries
- Inspected database schema and foreign-key relationships dynamically across 100,000 OpenPowerlifting competition results
- Added AST-based read-only SQL validation, automatic one-shot query repair using database error feedback, and grounded natural-language explanations of returned results