

Vladislav Fedorov

+1 (217) 974-9808 | vldslv.valeryevich@gmail.com

GitHub: github.com/ValDevelopment | LinkedIn: <https://www.linkedin.com/in/fedorov-vladislav-v>

Work Authorization: 36 Months STEM OPT (No Visa Sponsorship Required) | Available Immediately | Open to Relocation

PROFILE

MS Statistics graduate with experience building predictive models, reproducible data pipelines, and analytical applications using Python and SQL. Skilled in classification, time-series forecasting, feature engineering, model validation, and large-scale data analysis. Seeking entry-level machine learning engineering and applied data science roles.

EDUCATION

University of Illinois Urbana-Champaign | Champaign, IL | Aug 2024 - May 2026

Master of Science in Statistics | GPA: 3.90

University of Illinois Urbana-Champaign | Champaign, IL | Aug 2020 - Dec 2023

Bachelor of Liberal Arts and Sciences in Statistics, Minor in Mathematics | GPA: 3.77; Dean's List, Fall 2023

SKILLS

Programming: Python (Pandas, NumPy, scikit-learn), SQL, R, SAS

Machine Learning: Regression, Classification, SVM, XGBoost, PCA, Clustering, Image Classification, Feature Engineering, Model Selection, Cross-Validation, Hyperparameter Tuning

Data & Tools: MySQL, Git, GitHub, Streamlit, Tableau, Power BI, Excel

Languages: Russian (Native), English (Fluent), Italian (Intermediate)

PROFESSIONAL EXPERIENCE

Statistical Consultant | Champaign, IL

University of Illinois Urbana-Champaign | Jan 2026 - May 2026

- Conducted exploratory analysis of 6M+ job records and 100+ operational variables in Python and R, identifying resource inefficiency as a high-impact area and defining the scope of the subsequent analysis
- Built reusable Python and R programs that automate data preparation, resource-efficiency calculations, and visualizations, enabling the client to reproduce and extend the analysis with future job data
- Quantified systematic resource over-requesting, finding that the median job ran for 1.2 minutes against a 240-minute request and that approximately 800M reserved CPU-hours went unused
- Compared requested versus actual resource usage, revealing 1.48% median memory efficiency, 91.5% of jobs using less than half their requested memory, and pronounced CPU underutilization among jobs requesting 8–32 cores
- Collaborated with teammates and client stakeholders to present findings and propose personalized usage reports benchmarking each user's resource efficiency against cluster-wide patterns; the client identified the concept as a strong direction for future development

PROJECTS

Competitive Esports Predictive Modeling

- Built logistic regression and XGBoost models using professional League of Legends match data, achieving 0.84 AUC on a chronologically held-out season
- Forecasted weekly game-decisiveness trends using ARIMA, ETS, and gradient boosting with rolling-origin backtesting
- Built a sequential Elo rating system across 26,000+ professional matches to compare regional team strength, tune the K-factor empirically, and quantify uncertainty through bootstrapped confidence intervals

Personalized Strength Programming Agent

- Built an AI-powered application using Python, Pandas, Streamlit, and the Groq API to translate natural-language input into personalized squat, bench press, and deadlift programs
- Built a deterministic engine that converts a 39-session training log into percentage-based prescriptions
- Designed a hybrid LLM architecture in which the language model handles input interpretation and program explanations while a rule-based backend controls all training prescriptions

FAERS Pharmacovigilance Signal Detection

- Analyzed 385,288 FDA FAERS reports, including 3,279 semaglutide cases, to identify potential adverse-event signals
- Built parallel disproportionality-analysis pipelines in Python and SAS, calculating PRR, ROR, and chi-square statistics with Haldane-Anscombe correction and reproducing identical results across both implementations
- Identified 15 adverse-event signals and evaluated them against FDA labeling and peer-reviewed literature, visualizing disproportionality metrics in an interactive Tableau dashboard

OpenPowerlifting Relational Database

- Built a reproducible Python preprocessing pipeline to clean raw competition data, generate entity identifiers, preserve missing and failed attempts, and export relational tables for database loading
- Designed a normalized MySQL database for 100,000 OpenPowerlifting competition results, separating lifter, meet, and result data into linked tables with primary and foreign key constraints
- Wrote SQL analyses using joins, conditional aggregation, CTEs, and window functions, then built interactive Power BI dashboards to analyze federation participation, performance differences, sanctioned meets, and lift-attempt success rates